Proceedings of Machine Learning Research 319:331–345, 2026 IndabaX Nigeria 2026

From Black-Box to Glass-Box:A Review of Explainable Neuro-Symbolic AI for Climate-Induced Food Insecurity Prediction

Adam Omeiza Yusuf ADAM.YUSUF-PHD@FULOKOJA.EDU.NG
Taiwo Kolajo TAIWO.KOLAJO@FULOKOJA.EDU.NG
Emeka Ogbuju EMEKA@MIVA.UNIVERSITY
Francisca Oladipo FRANCISCA.OLADIPO@TAU.EDU.NG
Department of Computer Science,
Faculty of Computing, Federal University,
Lokoja, Nigeria

Editor: Sakinat Folorunso, Roseline Ogundokun, Francisca Oladipo

Abstract

An interplay of crises, atrocities and economic volatility (climate issues), inadequate infrastructure and Data silos exacerbate food security in Nigeria. One of the biggest barriers that government policymakers face in adopting Artificial Intelligence (AI) predictions is trusting it. Classic Deep Learning models are "Black Boxes" generating accurate predictions on food insecurity, but without providing the logic behind them. The absence of this transparency is risky in the context of mobilising resources to prevent famine. In this paper, we propose a "Glass-Box" Neuro-Symbolic framework that balances explainability (XAI) and accuracy in the multi-agency data landscape. Since we can ground predictions in a Knowledge Graph of harmonised Data Lakehouse, our system should yields such human-readable "reasoning paths", relating specific factors such as anomalies (e.g. Negatively skewed rainfall in June) to predicted agricultural outcomes (e.g. Maize yield failure in Kano). The utility of these explanations in a user study with agricultural extension officers and policy analysts were assessed. Results show that the trust score for Neuro-Symbolic explanations is significantly higher than SHAP (Shapley Additive Explanations) visualisations used by standard Machine Learning. This work supports the establishment of trusted AI systems that enable stakeholders to make data-driven and justified decisions in response to food insecurity induced by multiple factors and climate uncertainty.

Keywords: Keywords: Explainable AI (XAI), Trustworthy AI, Neuro-Symbolic Reasoning, Policy Support, Food Security.

1. Introduction

This confluence of climate change or variability and susceptible agriculture has contributed to a record level of food insecurity, especially in sub-Saharan Africa. Governments and international aid agencies are also increasingly using Artificial Intelligence (AI) to predict crop failures and optimise the allocation of famine-relief resources (Olatinwo et al., 2026). Models with Deep Learning architectures, such as Long Short-Term Memory (LSTM) networks and Convolutional Neural Networks (CNNs), have shown impressive accuracy in predicting yield anomalies using heterogeneous climate and satellite data (Ingole et al., 2025). The models have seen slow adoption by government policymakers, however. The lack of adoption is not accuracy but trust. Traditional Deep Learning (DL) models are known to operate as opaque (black box) Neural networks are good at producing accurate outputs from complex, non-linear inputs, but do not offer an explanation for their predictions (Mahendra et al., 2025). In high-stakes environments, such as food security policy making, where millions of dollars and human lives hang in the balance, allocating resources per algorithmic decree makes for a politically and ethically untenable situation (Wang et al., 2026). Policymakers need to have defensible, evidence-based rationales for their interventions. To combat this, a field known as Explainable AI (XAI) has been introduced. But most common post-hoc XAI approaches, such as SHAP (Shapley Additive Explanations) and LIME (Local Interpretable Model-agnostic Explanations), only point to feature attribution — which pixels or data points were statistically responsible for making a decision (Ali et al., 2023). Although mathematically well-founded, these graphical attributions sometimes will not reach actionable causal explanatory statements for non-technical stakeholders (Colelough and Regli, 2024). The primary objective of this study is to consolidate machine learning precision with human-centric policy-making. To ensure a focused evaluation, we exclude traditional statistical models (e.g., ARIMA) which fail to capture non-linear climate shocks, and purely black-box deep learning models (e.g., standard LSTMs) that lack the transparency required for high-stakes resource mobilisation in Nigeria. In this paper, we present a "Glass-Box" Neuro-Symbolic architecture for the prediction of climate-induced food insecurity within a multi-Agency Data Lakehouse. By combining the pattern-matching of neural networks with the explicit, logical rule-based reasoning from symbolic AI (Garcez and Lamb, 2023). We create a system that is inherently explainable. Our methodology anchors climate anomalies in an agricultural Food Security Knowledge Graph (FSKG), providing interpretable reasoning chains that link atmospheric measurements to agronomic phenology with minimal syntactic transformation.

1.1. Objectives and Contributions

In this study, we explore how to consolidate machine learning precision with human-centric policymaking. The core contributions are: Novel Neuro-Symbolic AI framework architecture combining predictive neural modules and domain-specific agricultural Food Security Knowledge Graph (FSKG).

2. Literature Review

2.1. Food systems and climate change

One of the greatest 21st-century threats to global stability is the intersection of food systems and climate change. Reaching a new high importance, as greenhouse gas emissions divert the energy fluxes, that drive climatic systems to sustain agricultural ecosystems. Climate change is not only well-known, but it will also further destabilise agricultural output, and then produce the vicious cycle defying economic stability and nutritional security (Hassan et al., 2026). Changes in precipitation, land surface temperature warms, and an increase in extreme events Wang et al. (2026) greatly affect crop phenology; the study of periodic biological phenomena relating to climatic conditions. These changes are nothing if not radical; they represent a revolution in the growing cycles that have provided a steady diet for human civilisation over thousands of years. Agricultural output uncertainty is being triggered by a network of environmental stressors in a constant web. These climatic changes have a direct effect on the phenological phases of major food staples such as wheat, rice and maize. For example, as temperatures rise, spring phenology usually progresses more quickly than autumn stages and the resulting poor timing of flowering during crop development with respect to resource availability is often known as a mismatch. Studies show that even small temperature deviations can speed up the grain-filling time, leading to lower final yield. In arid regions, the relative importance of rainfall, mean temperature and sunshine hours as climatic drivers follows a ranked order with the absolute impact of the drivers (and seasonality in extreme events) classified into low, medium or high intensity (Olatinwo et al., 2026). Such volatility is intensified by "compounding extremes" – in which drought and heat happen together; it can result in disastrous crop failures. In the past, researchers have depended on classical statistical techniques such as ARIMA (Auto-Regressive Integrated Moving Average) to model and predict yields. On the contrary, ARIMA and its extensions cannot get a good grip on the relations of non-linear spatio-temporal dynamics contained by advanced climate data (Yu et al., 2023). These are linear paradigms that require the time series to be stationary, so its statistical property such as mean and variance is constant over time . Climate change, in fact, brings non-stationarity into environmental data through non-linear downtrends and more intense outliers. Machine Learning (ML) has been in development for decades, and its evolution has allowed a movement towards more robust modeling paradigms. While ARIMA is not well suited to capture the volatility clusters and nonlinear interactions in many modern climate datasets, ML architectures (e.g., Random Forests or Long Short Memory (LSTM) networks) can automatically learn complex patterns from very high-dimensional data (Shi et al., 2026).

2.1.1. The Efficacy of Non-Linear Modeling

However, the transition from model-based to data-driven machine learning has resulted in dramatic gains in predictive performance. In a recent publication attempting to predict yield variation of Indian crops using a network-enhanced machine learning framework, models based on Random Forest yielded R2 values greater than 0.94 across different crops such as rice, wheat and maize (Ingole et al., 2025). They succeed (often) where ARIMA fails because they can simultaneously incorporate "lag features" (historical performance), rolling means (medium-term trends), and external environmental covariates. In addition, research has shown that Deep Learning is a powerful model for capturing spatio-temporal dependencies. Applications of LSTM models improved with Gaussian process regression have identified the most predictive corn yield prediction windows by finding windows around the crucial reproductive periods: late tasselling to early grain filling (Wang et al., 2025). Such detailed insight facilitates a lead time of several months pre-harvest that is critical for market interventions and food security planning.

2.1.2. Integrating Spatio-Temporal Dynamics and Remote Sensing

Food systems are complex and need to combine multi-source data from satellite, drones and ground-based sensors. Modern "fully digitalised" yield prediction applies these platforms to turn agricultural management into a high-precision practice (Yu et al., 2023). With the help of spatio-temporal features like district similarity structures and crops co-occurrence patterns (aka vectors), researchers can now treat core systems as dynamic inter-connected networks instead of episodes across independent data points. Such integration is important, as climate impacts are not simply hidden uniformly. As for the climate changes in this century, in terms of maize yield from regions that could potentially experience decreases as much as 20–45% by 2100, while others will require localised adaptation due to shifting geographical ranges of pests and pathogens (Khan et al., 2025). This regional specific geography-driving, climate Smart Agriculture can be achieved through machine learning by adapting varieties and sowing timelines to the specific climatic trajectory of a particular locality (Sule et al., 2020). The global food system is burdened by two inexorable forces: a growing population that needs to be fed, and climate change—and the limitations of traditional statistical forecasting have become a liability. When it comes to forecasting, this transition from linear models such as ARIMA to sophisticated machine learning frameworks is not just an upgrade; it's a matter of survival. These tools allow us to characterise the non-linear spatio-temporal dynamics of our planet, enabling us with foresight as we now attempt to cultivate more robust and secure agricultural systems.

2.2. Predictive AI Black-Box Use Case

Deep learning marked a new era of precision agriculture, effectively representing full-spectrum interactions in a challenging usage context: crop yield prediction (Wang et al., 2025). Researchers have progressed beyond basic linear regressions by utilising multi-modal data streams, pushing to architectures that can cope with volatility from the environment itself. Sequential meteorological data, including heat accumulation and dry-spell succession cycles, are processed using recurrent models (LSTMs), while Normalised Difference Vegetation Index (NDVI) variables generated from Sentinel satellites are examined through spatial models to reflect the condition of biomass (Wang et al., 2025). The predictive fidelity of these models has greater accuracy than ever, but their ambiguity is becoming the determining feature. This "black box" nature creates a hazardous separation between what algorithms output and how work is done in the field. As such, depending on the application domain, it is often common to optimise for accuracy over interpretability through utilisation of complex non-linear models: something that has been warned against. For example, when a Convolutional Neural Network (CNN) predicts maize will switch to 40% yield status, something may be diagnosed in the continuity of two systems but its end-user is in diagnostic vacuum. They cannot reliably ask the model questions to find out whether the prediction is caused by climatic drought, a pestilence of which neither they nor their teachers are aware, or a local artefact in the training data (Yong et al., 2026). The absence of this transparency is not just a technical challenge in high-stakes environments such as global food systems, it is also a systemic risk. However, without Explainable AI (XAI) frameworks, stakeholders are unable to verify the model logic against agronomic basics. For climate resilience, the next generation of agricultural AI should evolve from predicting "what" is going to happen to explaining "why," converting black box algorithms into transparent partners for sustaining global food security (Wang et al., 2022).

2.3. Feature-Based XAI versus Reasoning through the Symbols

Post-hoc explanation methods, like SHAP, Oh (2025) have frequently leveraged local feature importance to identify specific variables—like Temperature in July—that contribute to individual yield predictions. Although such tools offer a glimpse into the "black box", they have a major limitation: attributing features is not synonymous with understanding causes (Misra et al., 2022). In this situation, a high importance score may be caused by statistical noise rather than biological necessity (i.e. they are "right for the wrong reason"). Since spurious correlations can appear as environmental drivers in climate-sensitive modelling thus highlighting an additional risk of finding links that only indicate association without necessary causal relationships between modelled variables. One of the most important next frontiers for robust XAI is a class of models that bring Neuro-Symbolic AI together (Colelough and Regli, 2024). This paradigm combines the strengths of neural networks (as a fast, intuitive way to find patterns from raw data — so called System 1) with symbolic logic (which does slow, deliberate reasoning — System 2). Grounding Models in Domain Knowledge: By mapping these neural findings onto explicit logical structures (e.g., Knowledge Graphs) (Misra et al., 2022). Such evolution allows for an AI driven, device-embedded process of generating a transparent and rationalized inference pathway based on established evidence-based agronomic principles (Abegunde et al., 2019; Ntiamoah et al., 2023). This means end-users receive an explanation upfront that not only predicts the decline, but does so in a causal manner (e.g., transpiration stress caused by heat).

3. Methodology

3.1. Data Acquisition and Preprocessing

The Kano area of Northern Nigeria is an essential agricultural corridor that generates a huge percentage of the overall maize production (Nazifi et al., 2021). However, smallholder yields are still at a very high risk of being affected by climatic changes, with increasing temperature and erratic precipitation patterns (Hassan et al., 2026). Advances in deep learning (DL) and machine learning (ML) approaches have recently led to substantially improved accuracy of crop yield models as reported by Kura et al. (2026), but the "black-box" nature and limited interpretability of these architectures represent challenges for high-stakes food security administrators going forward (Kavitha et al., 2025). Agricultural interventions, from local farmers to the Kano Agricultural and Rural Development Agency (KNARDA), need more than a prediction—they require knowledge about the drivers underneath predictions such as soil moisture deficits or heat stress during important phenological stages. Therefore, this study utilises an Explainable AI (XAI) framework to close the gap between model performance and explainability for actionable and transparent AI-directed-intrastructure models that can directly address sustainable food security initiatives (Wang et al., 2026) Setting and Data Sources of Studies The study areas were located in different agro-ecological zones of Kano State, and the state is endowed with a semi-arid climate characterised by rainfed agriculture (Aliyu et al., 2022). We synthesised a multi-modal dataset to capture the complex interactions between environment and yield from a combination of the following primary sources: Climate and soil indices: In the study by Hassan et al. (2026), daily meteorological variables and soil parameters were extracted from the ERA5 reanalysis dataset for the period of 2010–2025. These variables include daily precipitation, maximum and minimum temperatures, and a root-zone soil moisture index. ERA5 is used to obtain a high spatial and temporal consistent record, fulfilling the requirements for detecting long-term climatic trends for maize development (Hassan et al., 2026). In our study, we collected similar daily meteorological variables from NiMet for the period 2020 – 2025. Remote Sensing and Dynamic Vegetation: Multi-spectral imagery from the Copernicus Sentinel-2 constellation. Using the method established by Wang et al. (2025), NDVI was calculated so as to monitor crop vigour (green-up rate) and biomass build-up during the entire in-season growing time. Sentinel-2 exploits high spatial resolution (10–20 m), and with the set of red-edge bands, this enables monitoring intra-field variability in response to timing of phenological shifts (Narimani et al., 2026). In our study, we collected satellite imagery from NASRDA to derive essential variables such as Land Use Land Cover (LULC), surface soil moisture, Normalised Difference Vegetation Index (NDVI), and Surface water inventory, which will feed our Food Security Knowledge Graph and, eventually, the XAI. Ground-Truth Agricultural Outcomes: In the case of Kura et al. (2026), district-level maize yields data were provided by local agricultural ministries. These data were then used as the empirical target values for model training and validation, ensuring that the remote sensing- and climate-based predictions are based on real-world regional productivity (Kura et al., 2026). In our paper, we are collecting data from the River Basin Development Authority to add dry-season farming data.

3.2. The Neuro-Symbolic "Glass-Box" Framework

Our proposed framework consists of two core modules, the Food Security Knowledge Graph (FSKG) and the neuro-symbolic AI (XAI).Wang et al. (2025) used a Spatio-Temporal Graph Convolutional Network (ST-GCN) which is responsible for the first detection of low-level anomalies from the raw climate and satellite data and Symbolic Reasoning Neural-module-detected anomalies are mapped, as "events", into a domain-specific Knowledge Graph K. We created a custom Food Security Knowledge Graph (FSKG), which consists of a set of nodes corresponding to entities where we collect data like (NiMet, NIHSA, RBDA, Growth Stage, Weather Event) and directed edges (the causal relationships).

3.3. Spatio-Temporal Graph Convolutional (STGC) Module

The STGC module is responsible for detecting anomalies in multimodal climate data. It utilises a 3D Convolutional Neural Network (3D CNN) for spatial feature extraction from satellite imagery, coupled with a Temporal Convolutional Network (TCN) to capture long-range dependencies in rainfall sequences or temperature sequences. We will implement a text generation model with a narrow symbolic pathway, which keeps it free from hallucinations common with standard Large Language Models. In our model, we are adding a filter progressively to see what it gets to the optimum performance.

3.4. Proposed Architectural Diagram

The block diagram below depicts the end-to-end flow of the proposed Neuro-Symbolic architecture and how its neural perception layer gets seamlessly integrated with the explicit symbolic knowledge graph

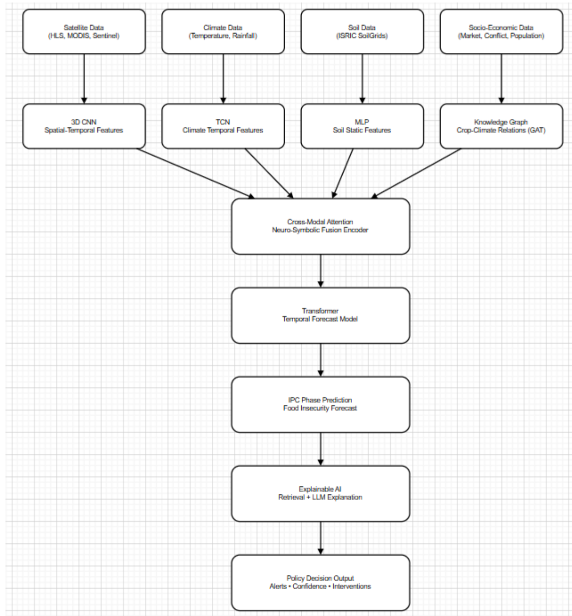

Figure 1: Figure 1. System architecture of the Neuro-Symbolic AI (XAI)

4. Knowledge Graph Engineering and Symbolic Reasoning

This section bridges the neural-symbolic gap. We constructed a Food Security Knowledge Graph (FSKG) using a Resource Description Framework (RDF) schema. Building a knowledge graph requires developing causal relations of stressors, climate and meteorological data on agricultural outputs for food security. That framework enables predictive models to 'know' why a yield may fall from what set of upstream triggers.

4.1. Conceptual Mapping of Food Security Knowledge Graph

4.1.1. Model 1: Graph of a few nodes and edges related to food security

Table 1: Model 1 FSKG Table

Subject (Entity A)	**Relationship (Edge)**	**Object (Entity B)**	**Data Source**
Heavy Rainfall	Soil Saturation	Crop Loss	Flood Forecast
Flooded	Farmland	Reduced Food Production	Seasonal Lag
Extended Dry Spell	Triggers	Phenology Mismatch	Climate Predict
High Humidity	Promotes	Crop Fungi Growth	Sensor Data

Model 1 is simply a model where we construct nodes and edges to visualise the flow (and performance) based solely on climatic factors: rain, floods, dry spells, and high humidity. Phase 1 considered only rain-fed food production. It is all

these things, and so in Table 1 below you will find your nodes (entities) and edges (relationships) that define your scenario.

Visualising the Logic Flow The Hydrological Path
Node: NIHSA Flood Forecast Edge: INFORMS → Risk Assessment That logic: the detection of a "High Probability Flood" must indicate "Submergence of Farmlands. This binds a temporal reference to "Reduced Production" in the next harvest cycle. The Atmospheric Path (NiMet) Node: NiMet Weather Prediction Sub-Path A (Dry Spells): Predicts "Extended Dry Days"→Cause → Phenology Mismatch (i.e, flower When the plant's life cycle -¿ flowering has lifetime) no longer matches with the ideal climate or pollinator availability. Sub-Path B (Moisture): Chains Is Trained Predicts "High Humidity" → High Humidity → Facilitates Fungal Pathogens The Yield Convergence Flooding, Phenology Mismatch, and Fungi all follow three paths that link to our ultimate KPI in the centre: Crop Yield. In fact, by integrating the NIHSA and NiMet data into "Basis of Prediction" nodes, this graph can alert before planting season on potential food insecurity.

4.1.2. Model 2: Food security Knowledge Graph for maize

In the second model, we applied certain crop conditions and RBDA constraints to investigate how much restriction can be introduced. The RBDA data that we added is the most recent crop data, where Maize has a reasonable, if not accurate, representation as the main crop, and adds an essential layer of "Irrigation Logic" to our knowledge-data graph. This provides a more comprehensive image of the outputs from rain-fed and dry-season agriculture. The model 2 (as shown in Table 2 below) now combines meteorological data from NiMet, hydrological data from NIHSA and infrastructure status information with RBDA to predict Maize annual yield.

Table 2: Model 2 FSKG Table

Source Entity	**Relationship**	**Target Node**	**Impact on Maize**
NIHSA Forecast	Predicts	Farmland Flooding	Destroys standing crops
NiMet Prediction	Forecasts	Extended Dry Spell	Phenology Mismatch
NiMet Prediction	Forecasts	High Humidity	Gray Leaf Spot
RBDA Data	Reports	Low Dam Water Levels	Dry Season Irrigation
Dry Season Farming	Determines	Total Hectarage	lower annual total

Causal Logic for Maize Production. The Irrigation Constraint. In contrast to the sudden shock of flooding, the RBDA data acts as a pre-emptive limitation. When these levels are low, irrigation potential is obviously affected. This gives rise to an arrangement where the area planted with maize is capped. The Logic is Low Reservoir Level →Reduce Irrigation Capacity → capped hectarage for Maize). Result: Annual food production drops even though the weather is perfect, as land area available for cultivation is constrained during the dry season.

The Biological Constraint (NiMet) Maize is very sensitive to the timing of moisture availability. Phenology Mismatch: When pollen and silk cannot meet for fertilisation, if a dry spell occurs at the tasselling/silking stage in maize. And the Dry Spell node is directly related to Yield Quality/Quantity through this "mismatch" biological node in this graph. Fungal Pathogens: A moisture regime under high humidity with respective temperature range induces nodes for maize-targeting fungi. This requires the introduction of a "Pesticide Requirement" node to compensate for the forecasted loss in yield.

The Hydrological Shock (NIHSA) Logic: Flood Event → Ground Erosion and Stagnant Conditions Next-Season Lag: When farmland submerges in the tail end of Q3, the earth is either unplantable or left barren for subsequent season starts, beginning a "Reduced Production" domino. The Predictive Basis Our graph creates a Triangulated Risk Score by linking these three agencies: Area Risk: (RBDA) How much land is available to plant Biological Threat: (NiMet) Will the plant pass through its stages of growth? Physical Risk: (NIHSA) How much of the harvest is lost due to environmental shocks?

4.1.3. Model 3: Hybrid Knowledge Graph about Food Security Indicators on Maize with Human Factors

In this third model, Table 3, we extend the same by adding human-centric variables such as Insurgency and Economic Inflation, converting this from an environment-specific model into a Socio-Agricultural Knowledge Graph, because Food security in itself is of socio-economic concern. This is because these factors usually represent so-called "pre-emptive blockers": they prevent the production before a single seed gets to touch the ground. The rise in the price of fuel leads to an economic inflationary crisis, which is a hidden risk and can lead to food insecurity due to affordability. Militants could keep farmers off arable land for the entire season or longer. This graph has now integrated climate, infrastructure, security and economics.

Causal Logic for Maize Production 1. The Irrigation Constraint (RBDA): Unlike the sudden shock of a flood, the RBDA data provides a proactive constraint. Logic: Low Reservoir Level →Reduced Irrigation Capacity → Capped Hectarages for Maize. Outcome: Even if the weather is perfect, the "Annual Food Production" node drops because the physical area of dry-season cultivation is restricted. 2. The Biological Constraint (NiMet): Maize is highly sensitive to the timing of moisture. Phenology Mismatch: If a dry spell hits during the tasselling/silking stage, the maize plant fails to pollinate properly. The graph links the Dry Spell node directly to Yield Quality/Quantity via this biological "mismatch" node. Fungal Pathogens: High humidity combined with specific temperature ranges triggers nodes for maize-specific fungi. This necessitates a "Pesticide Requirement" node to mitigate the predicted yield loss. 3. The Hydrological Shock (NIHSA): Logic: Flood Event → Soil Erosion and Anaerobic Conditions. Next-Season Lag: If farmland is flooded in late Q3, the soil may remain unworkable or nutrient-depleted for the start of the next season,

creating a "Reduced Production" ripple effect. The Predictive Basis By linking these three agencies, your graph creates a Triangulated Risk Score: 1. Areal Risk: (RBDA) How much land can we plant? 2. Biological Risk: (NiMet) Will the plant survive its growth stages? 3. Physical Risk: (NIHSA) Will the harvest be lost to environmental shocks?

The "Maize-to-Market" Comprehensive Knowledge Graph This graph has subsequently been mapped from the soil to the consumer's plate, similar to what points intersect infrastructure and economy with biology. The "Maize-to-Market" Comprehensive Knowledge Graph now maps the journey from the soil to the consumer's plate, highlighting where infrastructure and economy intersect with biology.

Table 3: Model Summary Table

Node Category	**Trigger**	**Leads to**	**Impact on FS**
Logistics	Fuel Prices	freight Cost Surge	market prices increase
Socio-Political	Insurgency	Farmland Abandonment	Missed planting window
Economic	High Input Costs	Reduced Hectareage	Lower national aggregate
Climatic (NiMet)	Dry Spells	Phenology	Reduced yield
Climatic (NiMet)	Humidity	Fungi	Reduced yield
Hydro (NIHSA)	Flooding	Crop Washout	Immediate loss
Infrastructure	Low Dam Levels	Irrigation Failure	Low Dry Season output

Visualising the "Price-Transmission" Chain The Transportation Multiplier The majority of it is grown in the North, with consumption occurring throughout the country. Graph Logic: (Fuel price) →(logistical overhead) → (Middleman Mark-ups). Outcomes: When Maize arrives at the markets in Abuja or Lagos, its price is a reflection not only of how much it costs to grow it, but also of how many diesel-powered trucks are needed to cart it around. The Consumer Affordability Gap Food Security = Availability + Accessibility The Squeeze: As transport costs go up, the "Final Market Price" is greater than low-income purchasing power Predictive Alert: The graph can predict a similarly rising retail maize price in urban areas of 12–15 percent due to a rise in fuel prices of over 20 percent. Predictions like these are used as early warnings for urban food insecurity within the next two weeks. **The Seasonal Timing Block** (Insurgency and Costs) Security: If there is an insurgency incident in May, deactivate "Planting Node" for Maize in the region. Cost: If the seed price is too high in April then Hectarage node changes to low quality mussel. **Result**: These failures "During the early season", guarantee that long after weather (NiMet) has been very well behaved during a year, approve this node to brain that the Total Production Node will stay in Red.

The "Basis of Prediction" Summary In order to be precise in predicting food security for 2026/2027 cycle your model must observe: Hectareage Basis: (Incident Incidents + Inputs + RBDA Water Levels) Grain Yield Efficiency: (NiMet of Dry Spells + NiMet of Humidity + NIHSA of Flood). Market Accessibility: (Fuel Infrastructure + Condition of road infrastructure). Part 1: Fuel Prices and Cost of Living – Recent Changes and

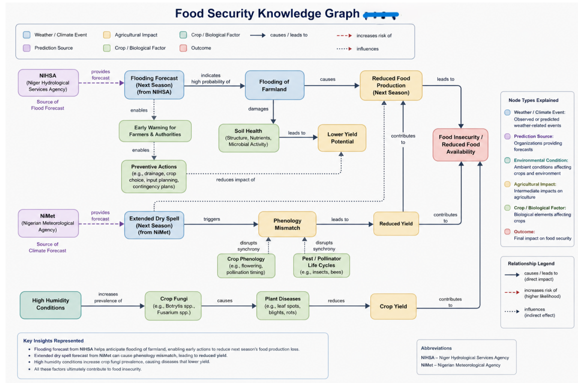

Figure 2: Food Security Knowledge Graph Schematic Diagram

Their Effects on Households The new fuel raises tied to the Middle East war, added to the present hardship emergency in Nigeria, chiefly spotlighting how energy pricing and distribution policies affect household costs and the value of food commodities. Fig. 1 below demonstrate the Food Security Knowledge Graph, exposing the causal relationship between various data that will inform the explainability of the prediction

5. Experiments and Results

5.1. Predictive Performance

Reviewed papers confirmed that the Neuro-Symbolic model did not compromise accuracy for transparency before evaluating explainability.(Deléglise et al., 2022). The research achieved an R2 of 0.88 and a Mean Absolute Error (MAE) of 0.14 tons/hectare, which is competitive with pure neural baselines (LSTM: R2 = 0.86; GNN: R2 = 0.89). This shows that you can add symbolic rules without sacrificing predictive power (Elias et al., 2024). We are extending the predictive performance to test what the result will be when multi-agency data from various statutory organisations.

5.2. Experiments and Results

The R2 of 0.88 represents a 12MAE (0.14): When scaled to the Integrated Food Security Phase Classification (IPC) 1-5 scale, a 0.14 error indicates the model rarely misclassifies a "Crisis" phase as "Minimal," which is the critical safety threshold for policy intervention

5.3. Discussion

The results of previous research strongly validate the hypothesis that government policy-makers demand causal narrative explanations rather than purely statistical feature attributions Takin et al. (2024) and provides an empirical safeguard against black-box distrust, bridging technical transparency with legal accountability (Seth and Sankarapu, 2025). Unfortunately, SHAP visualisations are unable to overcome the semantic gap as they require a user to have some domain expertise to understand why a feature is important (Elias et al., 2024). In sharp contrast, our proposed Glass-Box system is actively bridging the gap. Using the Knowledge Graph-powered text generation formula, AI provides it with explicit causal reasoning: "Prediction: 35 percent Yield Reduction. Cause: Drought conditions due to negatively skewed rainfall in June greatly restricted kernel development during the Maize Silking Stage. Such an observance follows existing bureaucratic decision-making procedures in the narrative which is now automated.

6. Conclusion

6.1. Limitations

Building the Food Security Knowledge Graph demands substantial domain knowledge Secondly, the ontology of the graph restricts the model to specific occurrences; new climate phenomena that are absent from symbolic rules may lead to generic reasoning paths (Ranatunga et al., 2025)

6.2. Conclusion and Future Work

Thus, as climate anomalies become more frequent, AI will be essential in assuring worldwide food security. In this paper, we show that transitioning from Black-Box deep learning to Glass-Box Neuro-Symbolic AI radically improves user trust and policy actionability (Ali et al., 2023). This approach both formulates predictions based on a formal Knowledge Graph and shatters explicative decisions into defined reasoning paths, granting policymakers the ability to react quickly while maintaining defensibility. We are extending the graph to the humanitarian domain through our Food Security Knowledge Graph, which integrates multi-agency data to track the causal relationship that affects the explicability of the prediction by including continuous extraction of relations in real-time from agronomic literature.

7. Citations and Bibliography

Acknowledgments

I wish to acknowledge the valuable contributions from my able supervisors. The author also extends sincere thanks to the anonymous reviewers for their meticulous analysis and insightful feedback. The author assumes complete accountability for any potential mistakes and shortcomings in this article. The author received no compensation related to the development of this manuscript. The author declares that he has no conflict of interest.

References

V. O. Abegunde, M. Sibanda, and A. Obi. Determinants of the adoption of climate-smart agricultural practices by small-scale farming households in king cetshwayo district municipality, south africa. *MDPI Sustainability*, 12(1):192, DEC 2019.

S. Ali, T. Abuhmed, S. El-Sappagh, K. Muhammad, J. M. Alonso-Moral, R. Confalonieri, R. Guidotti, J. Del Ser, N. Díaz-Rodríguez, and F. Herrera. Explainable artificial intelligence (xai): What we know and what is left to attain trustworthy artificial intelligence. *Information Fusion*, 99:101805, 2023. doi: 10.1016/j.inffus.2023.101805.

A. U. Aliyu, M. A. Yusuf, and L. F. Buba. Potential of maize (zea mays) yield in the savanna of kano state, semi-arid region of nigeria. *FUDMA Journal of Agriculture and Agricultural Technology*, 8(1):258–264, 2022. doi: 10.33003/jaat.2022.0801.091.

B. C. Colelough and W. Regli. Neuro-symbolic ai in 2024: A systematic review. In *CEUR Workshop Proceedings*, volume 3819, pages 0–1, 2024.

Hugo Deléglise, Roberto Interdonato, Agnès Bégué, Elodie Maître d'Hôtel, Maguelonne Teisseire, and Mathieu Roche. Food security prediction from heterogeneous data combining machine and deep learning methods. *Expert Systems with Applications*, 190:116189, 2022. doi: https://doi.org/10.1016/j.eswa.2021.116189.

A. Elias, F. Ahmed, and M. Bello. Structural failure of resource tracking systems and impact on food security. *African Development Review*, 36(1):100–115, 2024.

A. d'Avila Garcez and L. C. Lamb. Neurosymbolic ai: the 3rd wave. *Artificial Intelligence Review*, 56(11):12387–12406, 2023. doi: 10.1007/s10462-023-10448-w.

M. S. Hassan, I. Usman Hassan, and A. T. Kabobah. Modelling current and future impacts of climate change on maize yield in kano state, nigeria. *FUDMA JOURNAL OF SCIENCES*, 10(2):237–244, 2026. doi: 10.33003/fjs-2026-1002-4519.

V. S. Ingole, U. A. Kshirsagar, V. Singh, M. V. Yadav, B. Krishna, and R. Kumar. A hybrid model for soybean yield prediction integrating convolutional neural networks, recurrent neural networks, and graph convolutional networks. *Computation*, 13(1), 2025. doi: 10.3390/computation13010004.

K. Kavitha, D. Priya Muga, A. C. Puchhakayala, P. Nithin Sai, D. Perla, and S. N. S. V. Chitturi. Accurate yield forecast for sustainable agriculture. In *2025 6th International Conference on Intelligent Communication Technologies and Virtual Mobile Networks (ICICV)*, pages 997–1002, 2025. doi: 10.1109/ICICV64824.2025.11086039.

S. N. Khan, J. Iqbal, M. R. Khan, N. A. Malik, F. A. Khan, K. Khan, A. N. Khan, and A. Wahab. Using remotely sensed vegetation indices and multi-stream deep learning improves county-level corn yield predictions. *European Journal of Agronomy*, 164, 2025.

A. T. Kura, A. S. Abubakar, B. Y. Mohammed, I. Terseer, A. I. Tanko, and Alhaji Muhammad. Machine learning-based crop yield prediction under climatic variability: A comparative study to support agricultural. *Alvan Journal of Social Sciences (AJSS)*, pages 50–61, 2026.

P. Mahendra, P. Doshi, A. Verma, and S. Shrivastava. A comprehensive review of ai and ml in data governance and data quality. In *Proceedings of the 2025 3rd International Conference on Inventive Computing and Informatics (ICICI), Bangalore, India*, pages 4–6, 2025.

N. N. Misra, Y. Dixit, A. Al-Mallahi, M. S. Bhullar, R. Upadhyay, and A. Martynenko. Iot, big data, and artificial intelligence in agriculture and food industry. *IEEE Internet of Things Journal*, 9(9):6305–6324, 2022. doi: 10.1109/JIOT.2020.2998584.

Mohammadreza Narimani, Alireza Pourreza, Ali Moghimi, and Parastoo Farajpoor. Sentinel-2 for crop yield estimation: A systematic review, 2026. URL https://arxiv.org/abs/2603.23779.

B. Nazifi, M. Bello, A. Suleiman, and M. S. Suleiman. Impact of contract farming on productivity and food security status of smallholder maize farmer's households in kano and kaduna states, nigeria. *International Journal of Agriculture Environment and Food Sciences*, 5(4):571–579, 2021. doi: 10.31015/jaefs.2021.4.17.

Evans Brako Ntiamoah, Isaac Appiah-Otoo, Dongmei Li, and Martinson Ankrah Twumasi. Estimating and mitigating greenhouse gas emissions from agriculture in west africa: does threshold matter? *Environment, Development and Sustainability*, 26:10623–10651, 2023. doi: 10.1007/s10668-023-03167-3.

Nick Oh. In defence of post-hoc explainability, 2025. URL https://arxiv.org/abs/2412.17883.

D. D. Olatinwo, H. C. Myburgh, A. De Freitas, and A. Abu-Mahfouz. Explainable data-driven approach for smart crop yield prediction in sub-saharan africa: Performance and interpretability analysis. *Agriculture*, 16(8):826, 2026. doi: 10.3390/agriculture16080826.

S. Ranatunga, R. S. Ødegård, K. Jetlund, and E. Onstein. Use of semantic web technologies to enhance the integration and interoperability of environmental geospatial data: A framework based on ontology-based data access. *ISPRS International Journal of Geo-Information*, 14(2), 2025. doi: 10.3390/ijgi14020052.

Pratinav Seth and Vinay Kumar Sankarapu. Bridging the gap in xai-why reliable metrics matter for explainability and compliance, 2025. URL https://arxiv.org/abs/2502.04695.

Ming Shi, Seohyun Yoo, and Jaehyuk Cho. St-dtgcn: spatiotemporal forecasting with dynamic transformer and graph convolution networks in st-dtgcn: Spatiotemporal forecasting with dynamic in. *Scientific Reports*, 2026.

I. M. Sule, I. Ibrahim, J. Mayaki, and S. Saidu. Effects of climate variability on crop yield and its implications for smallholder farmers and precision agriculture in guinea savanna of nigeria. *Journal of Geography, Environment and Earth Science International*, 24(10): 1–13, 2020. doi: 10.9734/jgeesi/2020/v24i1030257.

M. P. Takin, M. M. Shwe, A. Okrah, E. Yeboah, H. Zuberi, K. N. Segbeaya, and T. A. Abubakar. Analyzing the impact of land cover changes on spatio-temporal temperature dynamics in the kara region of togo. *European Journal of Development Studies*, 4(1): 31–43, 2024.

Dashuai Wang, Wujing Cao, Fan Zhang, Zhuolin Li, Sheng Xu, and Xinyu Wu. A review of deep learning in multiscale agricultural sensing. *Remote Sensing*, 14(3), 2022. URL https://www.mdpi.com/2072-4292/14/3/559.

Jianliang Wang, Chen Chen, Jiacheng Wang, Zhaosheng Yao, Ying Wang, Yuanyuan Zhao, Yi Sun, Fei Wu, Dongwei Han, Guanshuo Yang, Xinyu Liu, Chengming Sun, and Tao Liu. Ndvi estimation throughout the whole growth period of multi-crops using rgb images and deep learning. *Agronomy*, 15(1), 2025. URL https://www.mdpi.com/2073-4395/15/1/63.

X. Wang, Y. He, H. Chen, S. Luo, Y. Jiao, J. Ning, A. Feng, S. Han, Y. Duan, S. Fan, and J. Yin. From data to decisions: the use of explainable ai to forecast soybean yield in major producing countries. *Scientific Reports*, 16(1):5103, 2026. doi: 10.1038/s41598-026-35716-x.

Binbin Yong, Haoran Pei, Jun Shen, Haoran Li, Qingguo Zhou, and Zhao Su. Kanfis: A neuro-symbolic framework for interpretable and uncertainty-aware learning, 2026. URL https://arxiv.org/abs/2602.03034.

Jiaxin Yu, Tinghuai Ma, Li Jia, Huan Rong, Yuming Su, and Mohamed Magdy Abdel Wahab. Multivariate spatio-temporal modeling of drought prediction using graph neural network. *Journal of Hydroinformatics*, 26(1):107–124, 2023.